

ChaosLemur: Injecting Failure into BGP

Davis Gossage
Duke University
dcg13@duke.edu

Emmanuel Shiferaw
Duke University
eas66@duke.edu

ABSTRACT

In this paper, we present ChaosLemur, a system built for applying Chaos, a recently developed and publicized principle in reliability engineering, to networks of virtual routers speaking BGP. The system is designed to take input from the user describing the desired number of routers, choice of topology, and distribution pattern of networks advertised in each router. Given this information, a set of linux containers are built, configured with software to act as routers, and connected in the described topology. The system then injects failures randomly into the network, collecting and saving the results/response of the system to those failures. We have implemented a prototype of ChaosLemur in Python, using Docker for containerization and Quagga as routing software and BGP implementation. As it currently stands, the system has been tested to configure and build networks of 5 - 20 containers in simple mesh and hub topologies speaking BGP, to take down specific and random nodes, and to report results in the form of logging routing tables.

1. INTRODUCTION

The first use of the word "Chaos" for the random, injected failure of elements in a distributed system was a Netflix technology blog post[4], which described a certain "Chaos Monkey", an actual unknown and inexplicable failure in Netflix's cloud infrastructure that drove engineers to focus more on fault-tolerance and reliability, both in the design and maintenance of the system. This phenomenon seems to be a simple result of the fact that failures that happen in production are hard to predict, and cannot be reliably recreated in traditional 'testing' environments. Thus, the injection of these failures frequently while in production acts as a sort of secondary testing phase, keeping engineers on their toes. This drove the team to build an automated version of this helpful bug that now runs in their production cloud on Amazon Web Services, taking down system elements at vastly different levels of scale, from single VM instances to entire AWS regions.[3]

Over time, Netflix has extended this idea to a more broad concept they have dubbed Chaos Engineering, and defined as the practice of designing and building experiments to understand the response of distributed systems to failure, and to inform action on reducing the negative effects of that response. This has been applied over the past few years to improve the reliability of Netflix, an extremely large and distributed cloud/web service that experiences many unpredictable failures in production. It has also been used at Microsoft (also for cloud services), as well as within the networking realm (specifically for SDN)[5].

As the design of any specific network topology/configuration is, in essence, the design of a distributed system, a system that would allow rapid simulation of many different configurations of a network may be quite useful. BGP as a whole has been criticized for its occasional reduced performance due to hot-potato routing/oscillations and long convergence time for knowledge of available paths[7]. It is also increasingly being used in data centers for the same reasons: since it does not require any "global" state or knowledge of networks, it can be used in scenarios where the overhead necessary for this global state (memory, network, etc.) would be too expensive.. Issues of oscillation and convergence can also be mitigated through tuning and proper configuration of a network [2]. The goal of ChaosLemur is to provide the ability to automate this tuning process and test a desired topology's fault-tolerance, without having to build it in hardware.

2. RELATED WORKS

The chaos principle developed by Netflix for the Cloud has been shown to be applicable to a significant and diverse array of distributed systems, including search [6].

In an attempt to consolidate efforts and formally define this chaos-driven approach to fault tolerance, Netflix has compiled the core ideas of the approach into a single document[1], emphasizing the importance of the 'Experiment' as a singular unit of information-discovery in the application of these principles. This manifesto puts forth a four-step process which simply revisits the basic scientific method, with the role of "Nature" being played by the unwieldy and unreliable distributed systems that power our most important large web services. The steps are as follows:

1. Start by defining 'steady state' as some mea-

surable output of a system that indicates normal behavior.

2. Hypothesize that this steady state will continue in both the control group and the experimental group.
3. Introduce variables that reflect real world events like servers that crash, hard drives that malfunction, network connections that are severed, etc.
4. Try to disprove the hypothesis by looking for a difference in steady state between the control group and the experimental group.

This analogy works well in that the complexity of distributed systems in the world cannot be easily built-up from first principles. The empirical approach attempts to build theories and models for systems out of small, testable hypotheses and experiments, which are the only tools available with such complex systems.

[9]:

3. ARCHITECTURE

ChaosLemur's architecture was structured primarily around the systems used to implement the components. We chose to use Docker containers instead of full virtual machines to implement each virtual router. Hardware-level virtualization was not necessary for our purposes; using a VMWare-like system would add unneeded overhead in the form of hypervisors and separate storage space for each virtual router [8]. Docker's interface for building and managing containers also lends itself well to scripting and automation.

Docker images, the base set of OS information/scripts used to build containers, can be generated from text configuration files called Dockerfiles. These can be easily maintained and distributed via popular version control software designed for tracking changes in text files. Dockerfiles describe the operating system, initial software to install, and files to load into the container from the host system.

Quagga is a software routing suite that implements many common protocols found in hardware routers, including BGP and OSPF. It has a simple architecture, with one daemon, *zebra*, acting as the central component of the system, interfacing directly with network hardware and presenting an API for communication. All protocol implementations are also individual daemons, each using the aforementioned *zebra* daemon's API to actually perform the routing. For the purposes of ChaosLemur, two of these additional daemons were necessary - *bgpd* and *vtys*. *bgpd* is the quagga daemon that implements the BGP protocol. It maintains the routing table/forwarding information base, and issues/receives announcements and withdrawals. *vtys* is the only Quagga daemon that does not implement a routing protocol; it presents a shell/command-line interface to *zebra* (and by proxy, the other daemons). Users interact with Quagga (issuing commands, retrieving information) using this CLI. *zebra*, *bgpd*

and *vtys* all need to be configured before running. Their configuration is loaded in text config files - one for each daemon (*zebra.conf*, *bgpd.conf*, *vtys.conf*). *Zebra* and *vtys* do not have a great deal of customizability in their configuration, but as would be expected, the protocol-implementation configuration files have to express quite a bit of information.

Topology information for the ChaosLemur must be communicated to each router through *bgpd.conf*. Each router must also have its own Dockerfile (in order to be built and run as a separate container). Due to these two constraints, the simplest way to generate an entire ChaosLemur experiment from a simple set of inputs was to build an entirely separate Docker 'build context' for each router in the network, loading in a distinct *bgpd.conf* file on each node. Each docker build context is a linux directory containing the dockerfile and all files referenced/needed to run the container. The portion of *bgpd.conf* that is generated by ChaosLemur is shown below:

```
router bgp 7675
bgp router-id 172.17.0.3
network 159.12.0.0/16
neighbor 172.17.0.2 remote-as 7675
neighbor 172.17.0.4 remote-as 7675
neighbor 172.17.0.5 remote-as 7675
```

The ChaosLemur API call/constructor that generated this specific configuration file portion was

```
CL = ChaosLemur(4, "hub", "pareto", 2.3).
```

This router is the second in the collection, with address 172.17.0.3, configured as the 'hub'. (Note: this API call generates all the routers/build contexts, not just that of this hub). The single "network" line is a result of the 'pareto' distribution input, which says that the number of networks advertised by this router should follow the pareto distribution, with an α value of 2.3.

4. FUTURE WORK

ChaosLemur as it currently exists is a decidedly simple prototype. Its simplicity proves the ease with which any organization or network operator could implement a similar solution for any real network that needs to improve its reliability. Much of the future work on this project would/should go toward solidifying it as a proof-of-concept, addressing edge cases and performing more thorough tests.

The configuration-file centric design of Docker and Quagga make ChaosLemur considerably extensible and flexible. As the "Context Generator" portion of the current implementation is essentially a program that writes *bgpd.conf* files given simple input, it could quite simply be extended to exploit any of the capabilities that Quagga's *bgpd* has that can be configured via *bgpd.conf*. The next option to be provided to users as input to ChaosLemur should be the ability to configure networks with multiple autonomous systems as in 1. BGP is by definition an external gateway protocol, and much of the useful information that could be collected




Figure 1: BGP topology with multiple Autonomous Systems.

from this system and applied in the real world regarding failures would have to do with convergence time amongst many different autonomous systems that are spread out. Another option that Quagga offers, which could be a further input to the system, is configuration of route maps (which would allow operators to configure routers to filter which networks are advertised).

Yet another, perhaps more interesting option would be to abstract ChaosLemur from *bgpd* and allow networks to be implemented that speak any protocol that Quagga implements. This, too, is made into a simple task by the fact that Quagga uses configuration files for all daemons with consistent structure and syntax. The exact same API methods that exist currently would work just as well for container-routers running *ospfd* or *olsrd*, the daemons for OSPF and OSLR respectively.

One way in which ChaosLemur's approach/application differs from all other "Chaos" implementations is the fact that it configures a 'test' environment itself, rather than causing actual failures in production. As BGP is a protocol implemented primarily in more expensive routers owned by ISPs and enterprises, it was not quite feasible to attempt to build a system that would work on production hardware and be simply developed using open source software-routing like Quagga. Quagga does, however, present the same abstractions as hardware routers, and a version of the ChaosLemur API could certainly be developed for a real ISP's BGP network. Other real-world Chaos applications have also taken the approach of configuring a separate, controlled sandbox separate from production to let failure-daemons run loose - the main benefit being that in automating the deployment of this environment, failures in the build/release/deploy process will be uncovered as well. This is perhaps less applicable for networks outside of SDN.

5. REFERENCES

- [1] Principles of Chaos Engineering. (English). *Principles of Chaos*.

- [2] P. L. E. N. Adel Abouchaev, Tim LaBerge. Brain-Slug: a BGP-Only SDN for Large-Scale Data-Centers. (English). *Microsoft Global Foundation Services Presentation*.
- [3] A. T. C. R. Ali Basiri, Lorin Hochstein. Chaos Engineering Upgraded. (English). *The Netflix Tech Blog*.
- [4] J. Atwood. Working with the Chaos Monkey. (English). *Coding Horror*.
- [5] T. Benson. Destroying networks for fun (and profit). (English). *HotNets*.
- [6] H. Nakama. Inside Azure Search: Chaos Engineering. (English). *Microsoft Azure Blog*.
- [7] J. R. Nick Feamster, Jay Borkenhagen. Guidelines for Interdomain Traffic Engineering. (English). *Computer Communication Review*.
- [8] R. R. J. R. Wes Felter, Alexandre Ferreira. An Updated Performance Comparison of Virtual Machines and Linux Containers. (English). *IBM Research Report*.
- [9] B. Wong. Introducing Chaos Engineering. (English). *The Netflix Tech Blog*.